

CprE 381: Computer Organization and Assembly-Level Programming

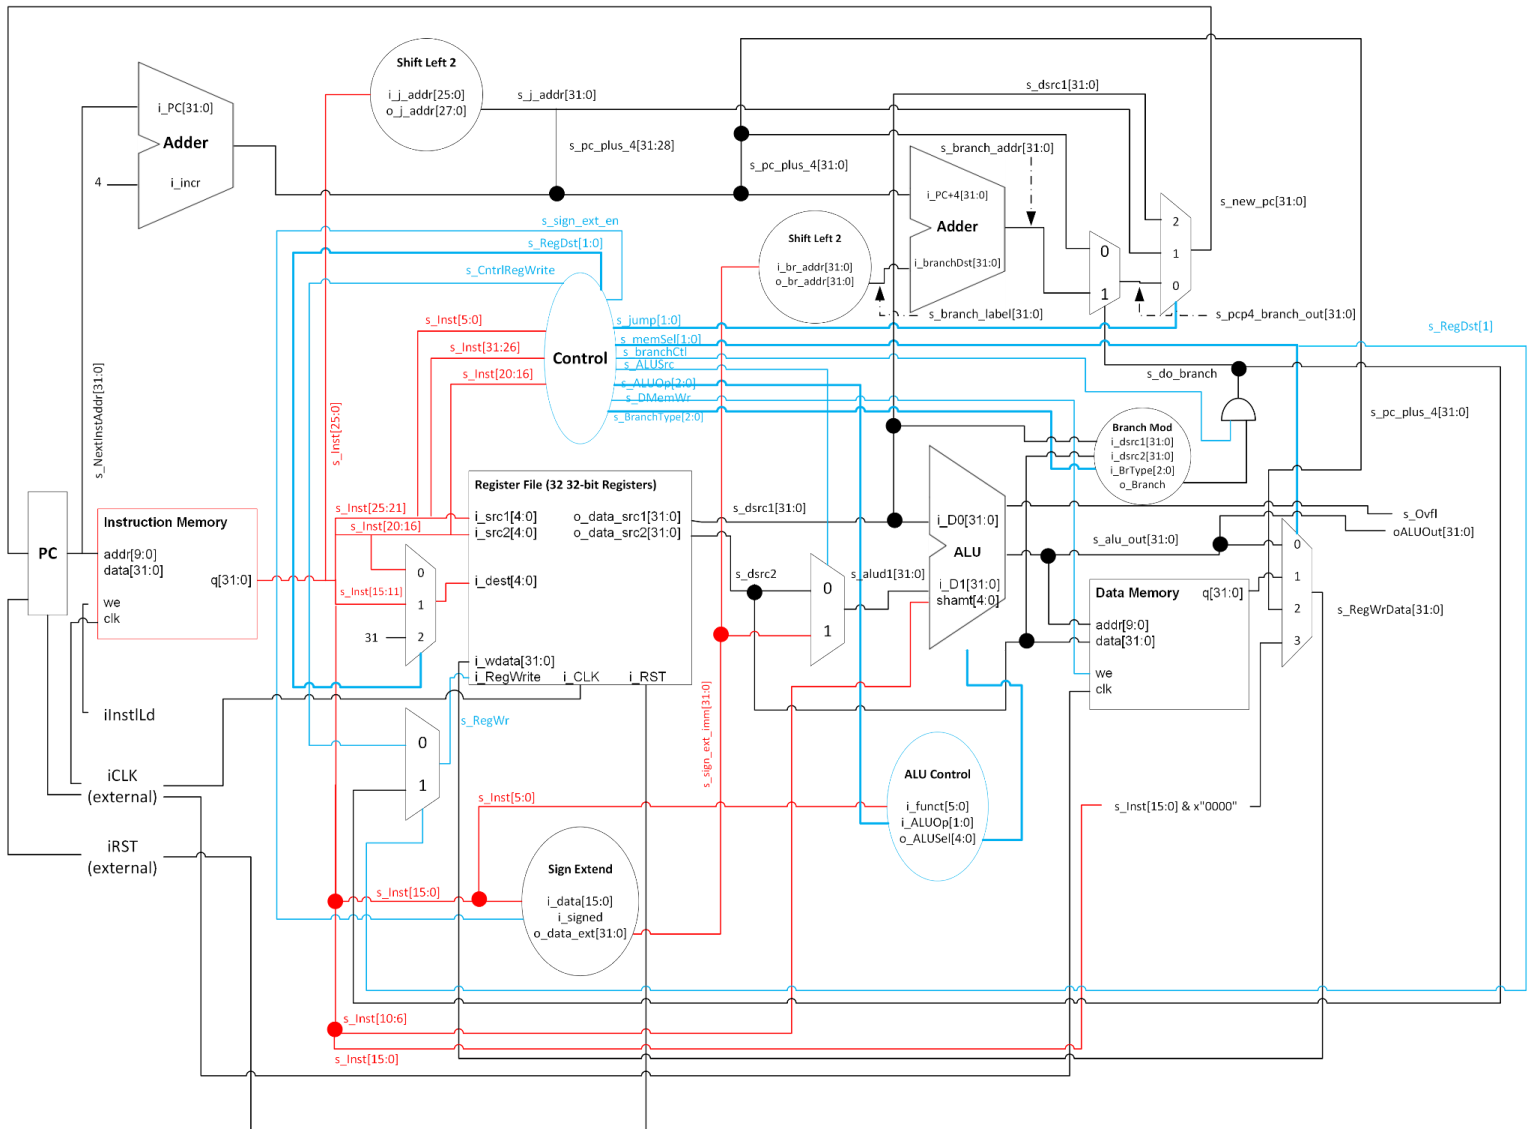
Project Part 1 Report

Team Members: Anthony Manschula
Henry Shires

Project Teams Group #: TermProj_3_02

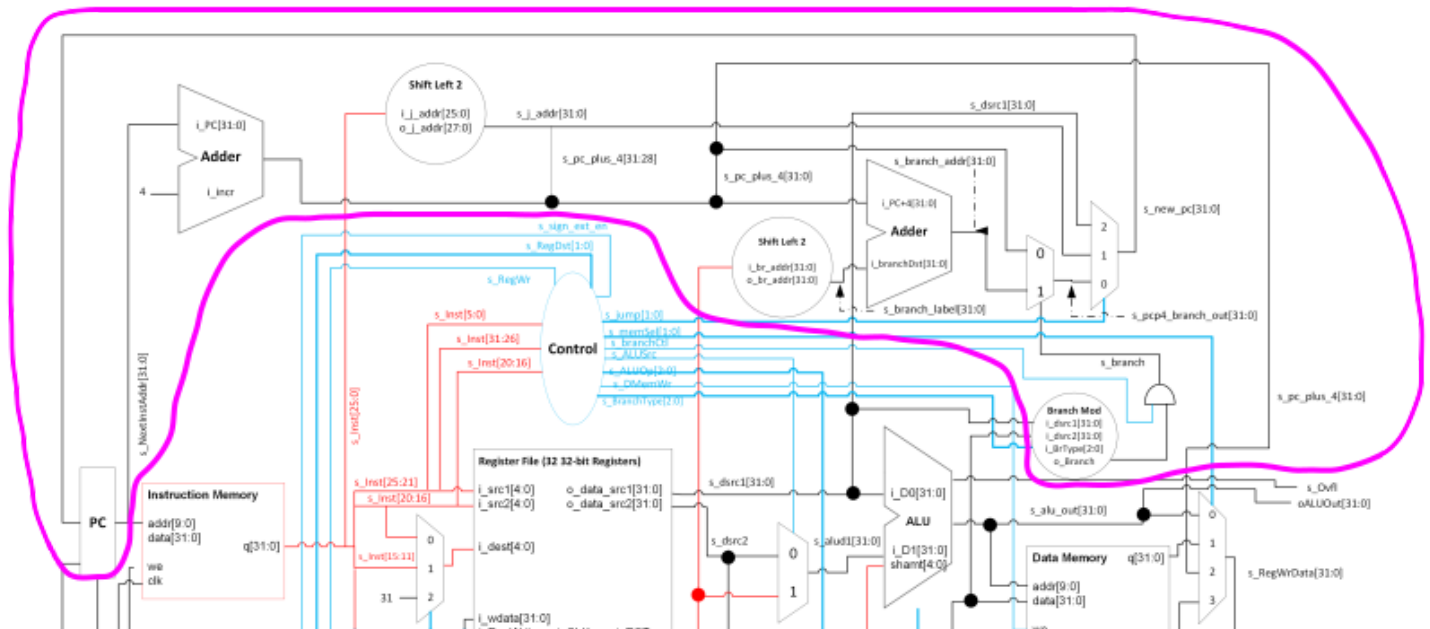
Please see [this folder](#) for full-resolution waveform images.

[Part 1 (d)] Include your final MIPS processor schematic in your lab report.



Our MIPS Single-Cycle Processor design



[Part 2 (b.ii)] Draw a schematic for the instruction fetch logic and any other datapath modifications needed for control flow instructions. What are additional control signals needed?



Additional control signals we elected to include in our design:

- s_jump : Controls whether the PC takes PC+4 (or branch address), a jump address specified by an immediate, or an address value specified in a register from jr .
- We also designed a standalone branch control module, which takes the $s_BranchType$ signal generated from the control unit and determines whether a branch should be taken based on input data and the type of control flow being performed.

[Part 2 (b.iii)] Implement your new instruction fetch logic using VHDL. Use QuestaSim to test your design thoroughly to ensure it works as expected. Describe how the execution of the control flow possibilities corresponds to the QuestaSim waveforms in your writeup.



- Sequential execution
 - The first 10 cycles are executing simple add, or, and, and lui instructions, as evidenced by s_Inst. It can be seen that the PC output data increments by 0x4 after the execution completes.
- Branching
 - The 11th cycle fetches a branch instruction. We can see that, although PC + 4 calculates to 0x0040002C, the input of the PC register is 0x00400030, which is the address of the branch. Moving down, we can ignore jump address since we are not jumping. We can see that the PC-relative branch address is calculated to be 0x4 (we are only branching over one instruction), and it is added PC+4 in order to obtain the final branch address. The branch control module compares the two inputs based on the i_BrType signal (0 =, 1 !=, 2 >= 0, 3 >0, 4 <= 0, 5 < 0) generated by the control module. In this case we are testing equality since BrType is 0, and since i_dsrc1 and i_dsrc2 are equal, the branch output is asserted. This same process applies for any of the other branch instructions. The output of the branch module controls the select line for the branch mux. Since it is 1, we choose to take the value of the branch address (calculated in section “Branch address”) instead of PC+4. We then move on to the final step, which is choosing the data to send back to the PC. The select line of this multiplexer also comes from the control unit, and in this case, a select value of 2b00 will result in the data from the branch mux (in the case of a branch this will be the branch address) being sent to the PC input via s_new_pc. There is also another case, which are the *b*al* instructions. These must set \$31 to the value of PC+4 when the branch is taken. If we jump to instruction 22 in the program (a BGEZAL), we can see that in addition to the aforementioned steps, we also set the select signal for the register write value mux to 2b10, which allows PC+4 to pass through to the write data port of the register file. A select signal also sets the destination register to a hardcoded decimal 31 (not shown).
- Jumping
 - The 14th instruction is a jump (at 280ns). We see that the input of the PC is 0x0040004C, which indicates that the jump address was properly calculated. Indeed, in the “Jump address” section, we can see that the input data from the instruction is properly shifted left two bits and resized. Skipping down to the “Final PC Mux” section, the calculated branch address is seen on the input bus, and the select signal coming from the control module is 1, which indicates we should send the jump address to the PC input. In the 15th cycle (300ns), we perform a *jal*. All of the relevant information just mentioned still applies, and we can see that we end up with a proper jump address calculated and sent to the PC. However, since we are linking, we also need to write PC+4 to \$31. We can see that the value of the register write value mux select signal is 2, which indicates we should pass PC+4 to the reg write port. The last jump case is *jr*, which is straightforward. The 25th instruction (500ns) is performing a *jr \$31* instruction to return from a *jal*. We simply set the select signal for

the Final PC Mux to 2b10, and this passes through the value of data source 1 from the register file.

[Part 2 (c.i.1)] Describe the difference between logical (*srl*) and arithmetic (*sra*) shifts. Why does MIPS not have a *sla* instruction?

Logical shifts do not account for the value of the bit at either extreme of the signal (i.e., bit 0 or bit 31) when shifting. They simply shift bits left or right and insert 0s at the opposite end of the direction of the shift. Arithmetic right shift preserves the sign bit's value by shifting the sign bit's value. MIPS does not have a *sla* instruction because the LSB doesn't represent much useful information in the context of an arithmetic shift.


[Part 2 (c.i.2)] In your write-up, briefly describe how your VHDL code implements both the arithmetic and logical shifting operations.

In logical shifting, bits left over after shifting are always set to 0. For the left shift, this means the bits to the right of the shifted bits are set to 0, and vice versa for the right shift. In arithmetic shifting, instead of replacing the leftover bits with 0s, they are replaced with the sign of the original binary number, in this case, using 2s complement, the most significant bit's value (MSB). Our VHDL implementation uses 2:1 multiplexers to manipulate each bit of the 32-bit input vector *i_shsrc*[31:0]. For each bit of the 5-bit shift amount vector *i_shamt*[4:0], a row of 32 2-1 MUXs with the corresponding shift amount bit as a selection for all is used to shift the bits X digits left or right. Suppose X is greater than the remaining bits available to shift. In that case, the MUX will use a 0 or 1 to assert on the shifted bit instead, with a 0 for logical operations and a 1 for arithmetic operations. The signal propagates through the X = 16-bit row, 8-bit row, 4-bit row, 2-bit row, and 1-bit row, a total of five rows, each matching a shift amount bit (1 or a 0). The outputted vector with the correct shift applied should appear at the end of the multiplexer array.

[Part 2 (c.i.3)] In your write-up, explain how the right barrel shifter above can be enhanced to support left-shifting operations.

Since the right barrel shift output bits directly manipulate *i_shsrc* in order, *i_shsrc* can be split up into bits 0-31 and inputted into the first row of 2-1 MUXs from left to right, shifting right based on the value of *i_shamt*. With left shifting, the current vector direction will result in an incorrect shift. While seemingly complicated, the right barrel shift can be easily modified to reverse the vector bits using an additional set of MUXs, where bit 31 of the original input is now located at bit 0 and vice versa, continuing until the entire vector has been flipped. This row of MUXs utilizes a separate select *i_shdir* where 0 corresponds to the left shift and 1 to the right shift. Each bit is either left alone or reversed with its opposite position (0 and 31, 1 and 30, etc.) based on the select value. Then the exact same right barrel shifter can be used with the reversed input for a left shift. The reverse vector process is repeated to undo the reversal of bits at the end of the barrel shifter to ensure the original input with the correct shift is outputted.

[Part 2 (c.i.4)] Describe how the execution of the different shifting operations corresponds to the QuestaSim waveforms in your writeup.



In our QuestaSim test bench, four common cases are tested from left to right. From 0-35, ns is sll with $i_shsrc = 0x00000001$, 35-70 ns srl with $i_shsrc = 0x80000000$, 70-105ns sra with $i_shsrc = 0x80000000$, and 105-140 ns sra with $i_shsrc = 0x08000000$. These specific input values help visualize the shift that occurs one after another. For cases 3 and 4, case 3 is run with the MSB of the input = 1 and case 4 with the MSB = 0 to ensure the arithmetic operation follows the sign of the vector. In all four cases, the '1' value on the input is seen shifted 1, 2, 4, 8, 16, and 31 bits as intended. Case

1 shifts it from LSB to MSB (left), case 2 from MSB to LSB (right), case 3 from MSB to LSB while tracing the sign of 1, and case 4 from MSB to LSB while tracing the sign of 0.

[Part 2 (c.ii.1)] In your write-up, briefly describe your design approach, including any resources you used to choose or implement the design. Include at least one design decision you had to make.

- MIPS PC
 - Design: We needed a component that was similar in behavior to a register but with a MIPS-compliant reset value. We followed the n-bit register design closely and included a reset value of 0x00400000.
- Shift Left 2
 - We needed a shift left 2 unit that could either take the form of an n-bit and n+2-bit input/output or an n-bit input/output. The choice was made to use a combination of dataflow and behavioral VHDL in order to create something that was reusable for both the jump address shift, as well as the branch address shift.
- Sign extender
 - We wanted an approach that was straightforward, so we opted to use a select statement in combination with some dataflow-style VHDL to map the input signal to the lower part of the output, and fill the upper part of the output with either 0's or the sign bit, depending the value of the sign extension enable control.
- ALU Control
 - The ALU control module needs to tell the ALU which arithmetic operation to perform based on the *funct* field of the instruction, as well as the ALUOp control output. The ALUSel output is selected for an R-type instruction when ALUOp is '000' and a matching *funct* is found on the input. The value of ALUOp may also be the same for some non-R-type instructions, such as `add` and `addi`.
- Branch Module
 - We quickly realized that the “zero” output of the ALU would not have been sufficient to support the required branch instructions and that trying to shoe-horn support into the ALU would have been messy and error-prone. We instead opted to add an entirely new module that outputs a “do branch” signal based on the two register outputs and a new control signal called BranchType, which is generated by the control and is unique to each type of branch instruction.


[Part 2 (c.ii.2)] Describe how the execution of the different operations corresponds to the QuestaSim waveforms in your writeup.

MIPS PC




- At its core, the MIPS PC is just a register with a custom reset value, and we've already validated the register design in previous labs. Therefore, we tested the reset functionality by making sure reset assigned the correct value, and that setting the input while reset is asserted did not affect the output in the first two clock cycles. From there, the register operates as expected with the reset signal not asserted.

Shift Left 2




- The shift left 2 module shifts the input left by two bits, and optionally increases the output width by two bits as well. Width increase is important because it is used in the jump address calculation part of the datapath. As it can be seen in the waveform, the input bits are shifted left two positions and the vacated positions are filled in with zeroes. Additionally, it can be seen that the output is resized correctly in the case of the top signal, and that it works properly for all 0's and all 1's input.

Sign Extender



- The sign extender simply maps the 16-bit input signal into the lower 16 bits of the output, and then concatenates either 0's or the value of input[15] onto output[31:16] when i_signed is 0 or 1, respectively. Due to the nature of this implementation, it will either work, or it will not work. There won't be some input values that it works for and some that it doesn't, so extensive testing would be redundant. It can be seen that for s_signed = 0, 0's are concatenated regardless of the value of input[15]. When s_signed = 1, the value of bit[15] is concatenated, as seen in the outputs for input 0x8000, as well as 0x1234 and 0x7F4C.

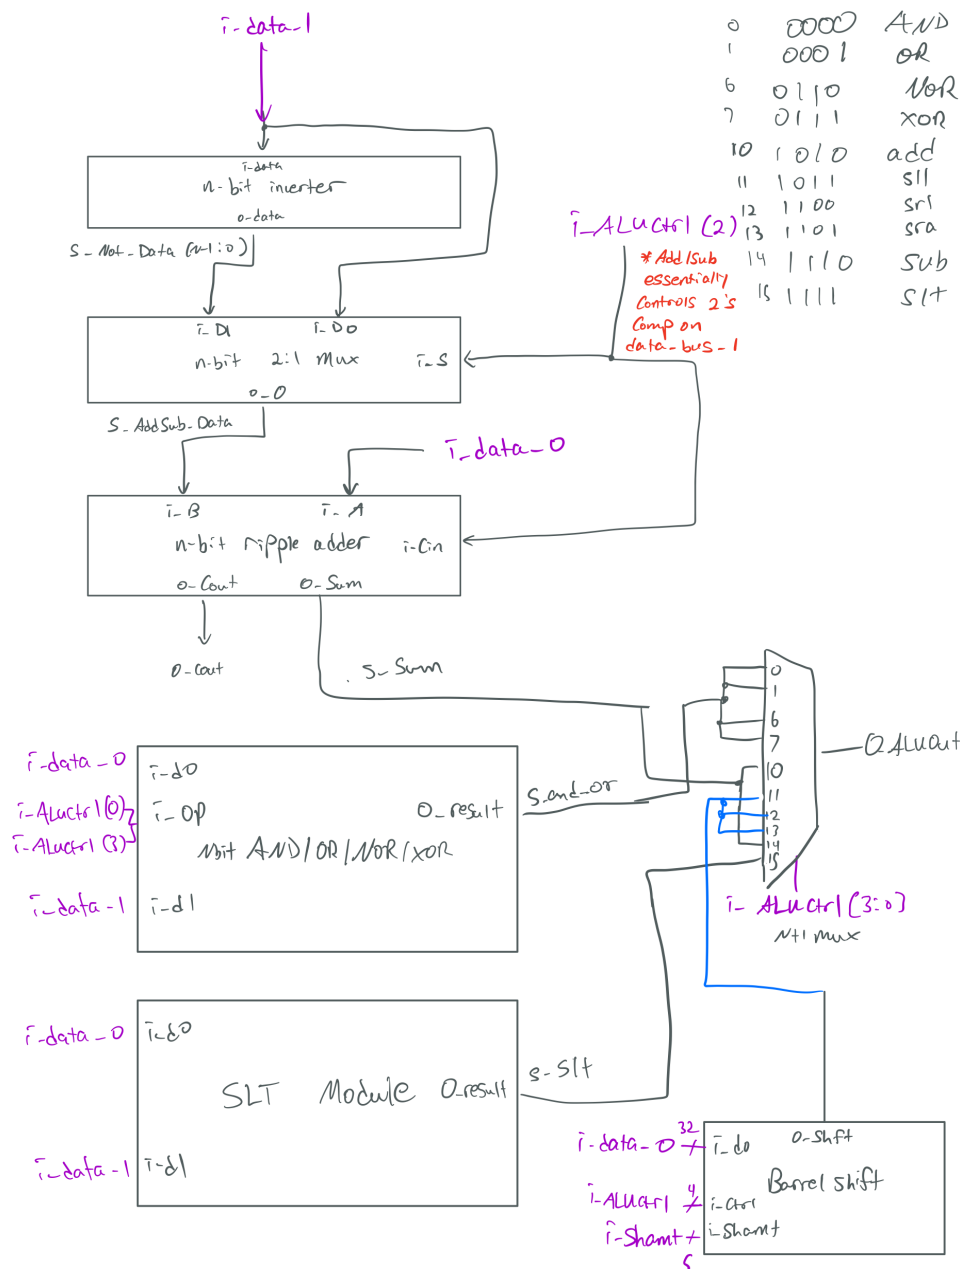
ALU Control



- ## Branch Module

- We test each branch type input with data that will cause the branch output to assert and data that will not cause the branch output to assert. We start with branch equal. The data is not the same, so it doesn't branch, but we see the output go high when it is. For a branch not equal (BrType = 1), it is the reverse of the previous condition. We test branch on greater than or equal to 0 (BrType = 2) for a negative, 0, and a positive, all giving the expected results. We follow that with a branch greater than 0 (BrType = 3) using a 0 input and a positive, also giving what we expect to see. We repeat the greater-than-equal test data, but instead, we choose a branch less than equal to 0 for the branch type (4), and the expected output is seen. The last branch type is less than 0 (5), and with 0 data, we expect to not branch, and with negative data, we do branch, which is what happens in reality.

[Part 2 (c.iii)] Draw a simplified, high-level schematic for the 32-bit ALU. Consider the following questions: how is Overflow calculated? How is Zero calculated? How is `slt` implemented?



Overflow occurs when the input signals have the same sign, and the output has the opposite sign when performing addition. When performing subtraction, overflow occurs if the signs of the inputs are different and the output has the same sign as the second data source. The overflow signal is only asserted when an instruction deals with data represented in a signed form. We do not calculate a zero signal since it was originally used for branching, but since we have a standalone branch module, it is redundant. SLT is

implemented via a simple behavioral module that compares the two (signed) inputs and asserts an output if the first input is less than the second.

[Part 2 (c.v)] Describe how the execution of the different operations corresponds to the QuestaSim waveforms in your writeup.


In our test bench, our ALU is operating as expected. Addition operations are either incrementing the number or wrapping it around if it reaches a result requiring more than 32 bits to represent. This is also the case for subtraction, but with decrementing the number. AND and OR also behave as expected, following the respective truth tables outputs.



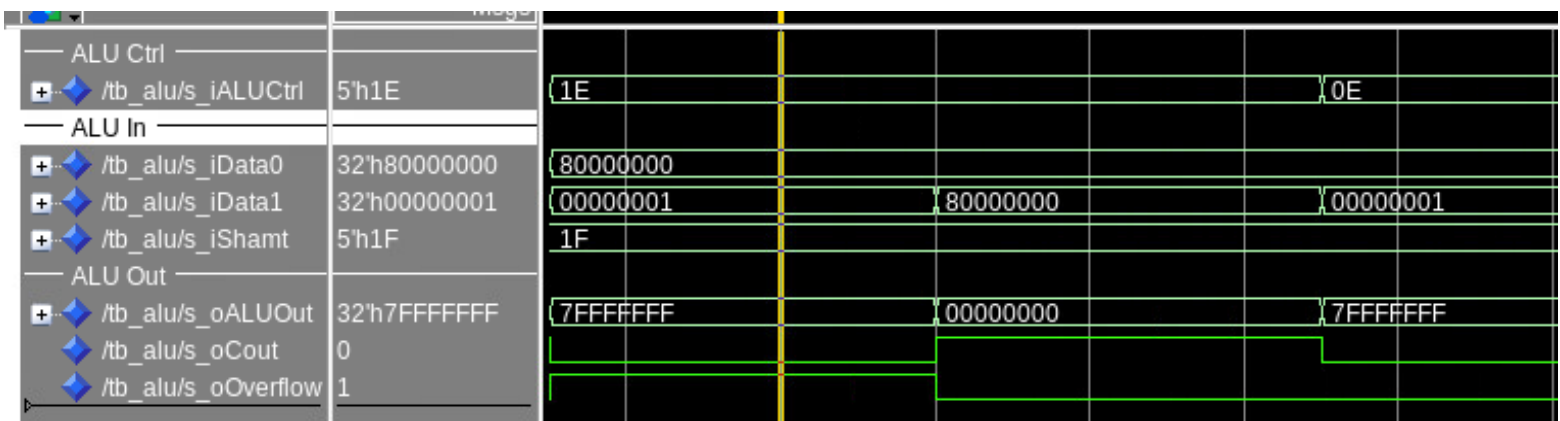
Add (signed with overflow then unsigned)



AND common cases from the truth table



OR common cases from the truth table + edge case



Sub (signed with overflow then unsigned)

[Part 2 (c.viii)] Justify why your test plan is comprehensive. Include waveforms that demonstrate your test program's functioning.

To ensure our testbench is comprehensive, we've included multiple common and edge cases for each instruction where applicable. For the addition operation, we've included a normal, signed case, a normal unsigned case, and an overflow signed case. This is the same for the subtraction operation. For AND, we incorporated a bit-addressable example and all 0s or all 1s in the vector and changed when they were used for four cases. We replicated the alternation of using 0x00000000 and 0xFFFFFFFF for AND and other logical operations (NOR, XOR, and OR). These cases matched the general idea of their truth tables but at a byte level, while our edge case addressed individual bits for each. For set less than operations, we included changes in the sign of each input, alternating between two positives, one positive and one negative (and vice versa), and both negatives. Finally, for sll, srl, and sra, each case has a single 1 on the input and shifts by the next bit (0, 1, 2, 4, 8, 16, and 31) during each succeeding case. This produces the same output from the barrel shifter test bench to ensure an easy match between the two outputs.

Overflow is asserting only on signed operations and asserting only when overflow occurs.



NOR

| | | |
|---------------------|--------------|--------------------------------------|
| | Msgs | |
| ALU Ctrl | | |
| /bb_alu/s_iALUCtrl | 5'h07 | 07 |
| ALU In | | |
| /bb_alu/s_iData0 | 32'h5152CA3E | 5152CA3E000000000000000000000000 |
| /bb_alu/s_iData1 | 32'h0F0F0F0F | 0F0F0F0F00000000FFFFFF00000000FFFFFF |
| /bb_alu/s_iShamt | 5'hXX | XX |
| ALU Out | | |
| /bb_alu/s_oALUOut | 32'h5E5DC531 | 5E5DC53100000000FFFFFF00000000 |
| /bb_alu/s_oCout | 1 | |
| /bb_alu/s_oOverflow | 0 | |

XOR

| | | |
|---------------------|--------------|--------------------------------------|
| ALU Ctrl | | |
| /bb_alu/s_iALUCtrl | 5'h0F | 0F |
| ALU In | | |
| /bb_alu/s_iData0 | 32'h70000000 | 70000000700012340000000000000000 |
| /bb_alu/s_iData1 | 32'h70001234 | 7000123470000000FFFFFF00000000FFFFFF |
| /bb_alu/s_iShamt | 5'hXX | XX |
| ALU Out | | |
| /bb_alu/s_oALUOut | 32'h00000001 | 00000001000000000000000100000000 |
| /bb_alu/s_oCout | 0 | |
| /bb_alu/s_oOverflow | 0 | |

SLT

| | | |
|---------------------|--------------|--|
| ALU Ctrl | | |
| /bb_alu/s_iALUCtrl | 5'h0F | 0B |
| ALU In | | |
| /bb_alu/s_iData0 | 32'h70000000 | FFFFFFF000000001 |
| /bb_alu/s_iData1 | 32'h70001234 | 00000001 |
| /bb_alu/s_iShamt | 5'hXX | 000101020408101F |
| ALU Out | | |
| /bb_alu/s_oALUOut | 32'h00000001 | 000000010000000200000004000000100000010000100080000000 |
| /bb_alu/s_oCout | 0 | |
| /bb_alu/s_oOverflow | 0 | |


SLL

| | | |
|---------------------|--------------|--|
| ALU Ctrl | | |
| /bb_alu/s_iALUCtrl | 5'h0C | 0C |
| ALU In | | |
| /bb_alu/s_iData0 | 32'hFFFFFFFE | FFFFFFF000000000 |
| /bb_alu/s_iData1 | 32'h80000000 | 00 |
| /bb_alu/s_iShamt | 5'h00 | 0001020408101F00 |
| ALU Out | | |
| /bb_alu/s_oALUOut | 32'h80000000 | 80000000400000002000000008000000008000000000000000000000000000 |
| /bb_alu/s_oCout | 1 | |
| /bb_alu/s_oOverflow | 0 | |

SRL

| | | |
|---------------------|--------------|--|
| ALU Ctrl | | |
| /bb_alu/s_iALUCtrl | 5'h0D | 0D |
| ALU In | | |
| /bb_alu/s_iData0 | 32'hFFFFFFFE | FFFFFFF000000000 |
| /bb_alu/s_iData1 | 32'h80000000 | 80000000 |
| /bb_alu/s_iShamt | 5'h01 | 01020408101F |
| ALU Out | | |
| /bb_alu/s_oALUOut | 32'hC0000000 | C0000000E0000000F8000000FF800000FFFF8000FFFFFF |
| /bb_alu/s_oCout | 1 | |
| /bb_alu/s_oOverflow | 0 | |

SRA



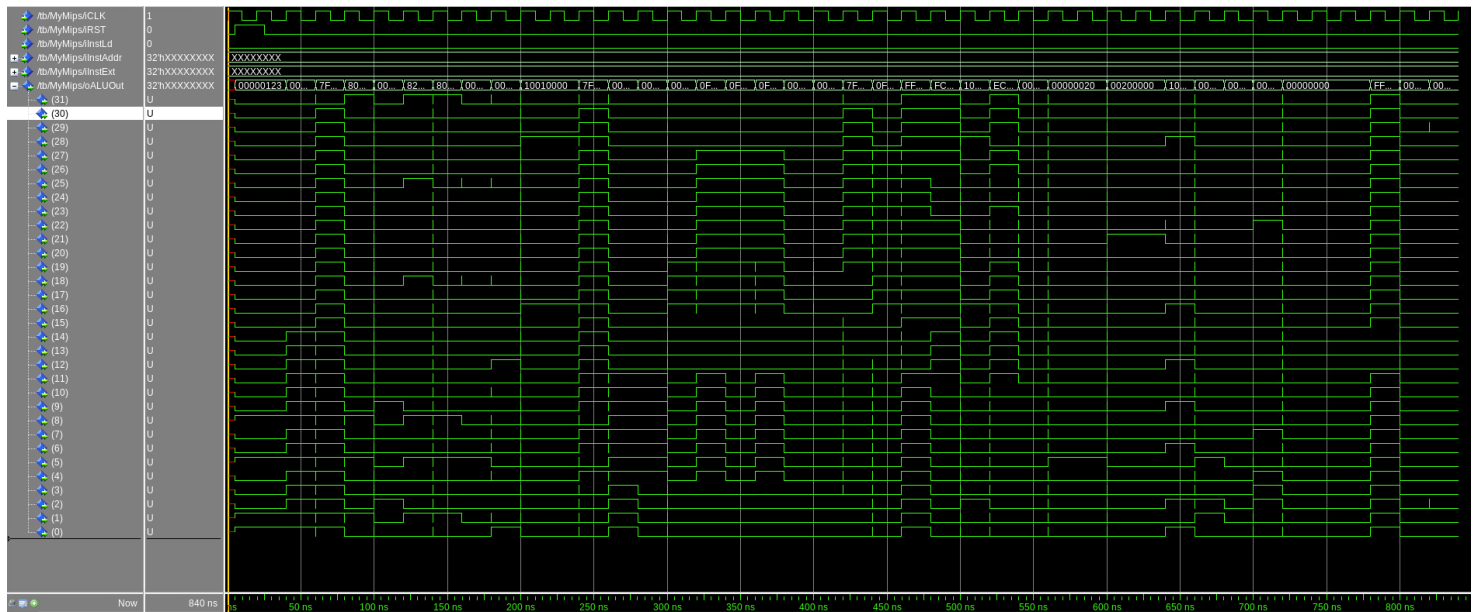
Reproduced Barrel Shifter Output (the arrow is the start of sll tests). The stack of waveforms matches the barrel shifter test bench earlier in this report.

[Part 3] In your writeup, show the QuestaSim output for each of the following tests and provide a discussion of result correctness. It may be helpful to also annotate the waveforms directly.


In all three waveforms, with each instruction usually resulted in a register or memory update. Each waveform shows the output of the ALU, which should correspond with any register changes. For example, in the base_test program, register \$8 is written \$t1 + 0x123, which is populated on the waveform from the ALU.

Additionally, in cf_test, \$4 is written 0 + 6, which populates on its waveform as 0x6, which is expected. Finally, the Bubblesort algorithm had a very large execution time and generated thousands of nanoseconds of signal propagations due to its recursive nature and the intensity of the sorting algorithm. After ~20,000 nanoseconds, Bubblesort completes and calls a halt, which sets the ALU output to 0 to exit the program seen on the waveform. Additionally, the address of the array data label can be seen computed after several previous instructions, which are used to reference the items to be sorted.


[Part 3 (a)] Create a test application that uses every required arithmetic/logical instruction at least once. The application need not perform any particularly useful task, but it should demonstrate the full functionality of the processor (e.g., sequences of many instructions executed sequentially, 1 per cycle, while data written into registers can be effectively retrieved and used by later instructions). Name this file Proj1_base_test.s.



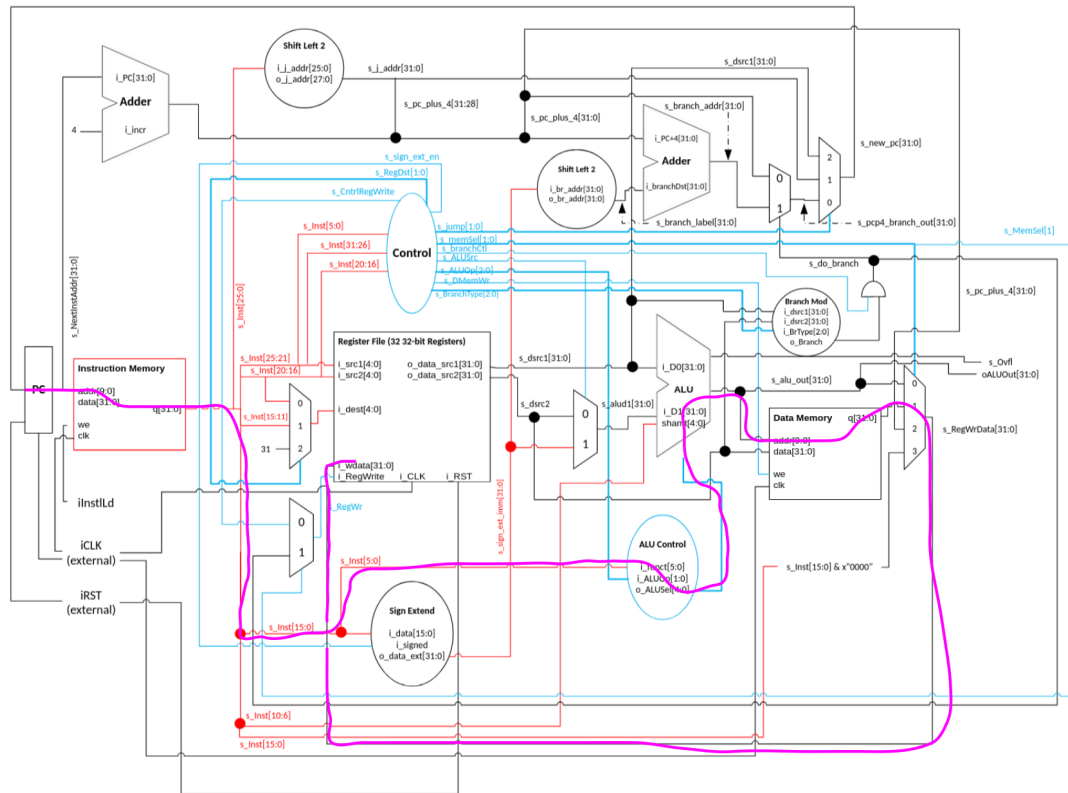
[Part 3 (b)] Create and test an application that uses each of the required control-flow instructions and has a call depth of at least 5 (i.e., the number of activation records on the stack is at least 4). Name this file Proj1_cf_test.s.



[Part 3 (c)] Create and test an application that sorts an array with N elements using the BubbleSort algorithm ([link](#)). Name this file Proj1_bubblesort.s



[Part 4] Report the maximum frequency your processor can run at and determine what your critical path is. Draw this critical path on top of your top-level schematic from part 1. What components would you focus on to improve the frequency?



- The max frequency is 22.49MHz.
- The critical path is PC > IMEM > ALU Control > ALU > DMEM > DMEM Multiplexer > Write to reg.
- It would probably make the most sense to optimize the ALU Control since there are 3.5ns of latency alone in the first step of the path through it. The barrel shifter also adds quite a bit of latency, so it might be interesting to see what a non-structural implementation could result in.